

# fourier\_transform\_draft

April 21, 2025

## 1 Volume 2: The Discrete Fourier Transform

```
[25]: import IPython
      from IPython.display import Audio, display
      import numpy as np
      from scipy.io import wavfile
      from scipy.fftpack import fft
      from matplotlib import pyplot as plt
      from scipy.io.wavfile import write

[26]: plt.rcParams["figure.dpi"] = 300          # Fix plot quality.
      plt.rcParams["figure.figsize"] = (12, 3)  # Change plot size / aspect (you
      ↪ may adjust this).

[27]: class SoundWave(object):
      """A class for working with digital audio signals."""

      # Problem 1
      def __init__(self, rate, samples):
          """Set the SoundWave class attributes.

          Parameters:
              rate (int): The sample rate of the sound.
              samples ((n,) ndarray): NumPy array of samples.
          """
          self.rate = rate
          self.samples = samples
          self.time = np.arange(len(samples)) / rate
          self.freq = np.fft.rfftfreq(len(samples), d=1/rate)

      # Problems 1 and 7
      def plot(self, show_dft=False):
          """Plot the graph of the sound wave (time versus amplitude)."""
          if show_dft:
              fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 6))

              # Time-domain plot
```

```

ax1.plot(self.time, self.samples)
ax1.set_xlabel("Time (s)")
ax1.set_ylabel("Amplitude")
ax1.set_title("Sound Wave")

# Frequency-domain plot
n = len(self.samples)
freq_data = fft(self.samples)
magnitude = np.abs(freq_data)[:n // 2]
freqs = np.arange(n // 2) * self.rate / n

ax2.plot(freqs, magnitude)
ax2.set_xlabel("Frequency (Hz)")
ax2.set_ylabel("Magnitude")
ax2.set_title("Frequency Spectrum (DFT)")

plt.tight_layout()

else:
    plt.plot(self.time, self.samples)
    plt.xlabel("Time (s)")
    plt.ylabel("Amplitude")
    plt.title("Sound Wave")
    plt.ylim([-32768, 32767])

# Problem 2
def export(self, filename: str, force=False) -> None:
    """Generate a wav file from the sample rate and samples.
    If the array of samples is not of type np.int16, scale it before
    ↪ exporting.

    Parameters:
        filename (str): The name of the wav file to export the sound to.
    """
    if self.samples.dtype == np.int16 and not force:
        write(filename, self.rate, self.samples)
        print(f"Exported without scaling: {filename}")
        return

    # Scale to int16 using formula (7.2)
    data = self.samples.astype(np.float64)
    peak = np.max(np.abs(data)) or 1.0 # Prevent divide-by-zero
    scaled = (data / peak * 32767).astype(np.int16)
    write(filename, self.rate, scaled)
    print(f"Exported with scaling: {filename}")

# Problem 4

```

```

def __add__(self, other):
    """Combine the samples from two SoundWave objects.

    Parameters:
        other (SoundWave): An object containing the samples to add
            to the samples contained in this object.

    Returns:
        (SoundWave): A new SoundWave instance with the combined samples.

    Raises:
        ValueError: if the two sample arrays are not the same length.
    """
    if not isinstance(other, SoundWave):
        return NotImplemented # allow NumPy to try
    other dtypes

    if len(self.samples) != len(other.samples):
        raise ValueError("SoundWave objects must have the same number of
    samples")

    # Promote to float64 to minimise clipping/overflow when summing
    combined = self.samples.astype(np.float64) + other.samples.astype(np.
    float64)
    return SoundWave(self.rate, combined)

# Problem 4
def __rshift__(self, other):
    """Concatenate the samples from two SoundWave objects.

    Parameters:
        other (SoundWave): An object containing the samples to concatenate
            to the samples contained in this object.

    Raises:
        ValueError: if the two sample rates are not equal.
    """
    if not isinstance(other, SoundWave):
        return NotImplemented # allow NumPy fallback

    if self.rate != other.rate:
        raise ValueError("Sample rates differ: {} Hz vs {} Hz"
            .format(self.rate, other.rate))

    concatenated = np.concatenate([self.samples, other.samples])
    return SoundWave(self.rate, concatenated)

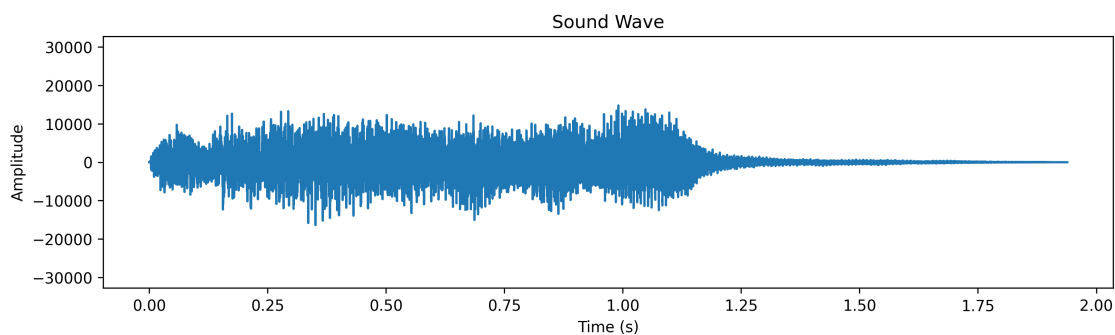
```

## 1.1 Problem 1

- Implement `SoundWave.__init__()`.
- Implement `SoundWave.plot()`.
- Use SciPy's `wavfile.read()` and the `SoundWave` class to plot `tada.wav` in the cell below.

```
[28]: rate, samples = wavfile.read("tada.wav")
sw = SoundWave(rate, samples)
try:
    sw.plot()
except NotImplementedError:
    pass

plt.show()
```



## 1.2 Problem 2

- Implement `SoundWave.export()`.
- Use the `export()` method to create two new files containing the same sound as `tada.wav`: one without scaling, and one with scaling (use `force=True`).
- Use `IPython.display.Audio()` to embed the original and two new versions of `tada.wav` in the cells below.

```
[29]: # (a) No scaling
sw.export("tada_noscale.wav")

# (b) With scaling - force=True
sw.export("tada_scaled.wav", force=True)
```

Exported without scaling: `tada_noscale.wav`  
Exported with scaling: `tada_scaled.wav`

```
[30]: display(Audio("tada.wav", autoplay=False, element_id='orig'))
display(Audio("tada_noscale.wav", autoplay=False, element_id='noscale'))
display(Audio("tada_scaled.wav", autoplay=False, element_id='scaled'))
```

<IPython.lib.display.Audio object>

<IPython.lib.display.Audio object>

<IPython.lib.display.Audio object>

### 1.3 Problem 3

- Implement `generate_note()`.
- Use `generate_note()` to create an A tone that lasts for two seconds. Embed it in the cell below.

```
[31]: def generate_note(frequency, duration):  
    """Generate an instance of the SoundWave class corresponding to  
    the desired soundwave. Uses sample rate of 44100 Hz.  
  
    Parameters:  
        frequency (float): The frequency of the desired sound.  
        duration (float): The length of the desired sound in seconds.  
  
    Returns:  
        sound (SoundWave): An instance of the SoundWave class.  
    """  
    sample_rate = 44100 # Standard sample rate for audio  
    t = np.linspace(0, duration, int(sample_rate * duration), endpoint=False)  
    samples = np.sin(2.0 * np.pi * frequency * t).astype(np.float32)  
    return SoundWave(sample_rate, samples)
```

```
[32]: a_tone = generate_note(440.00, 2)  
b_tone = generate_note(493.88, 2)  
c_tone = generate_note(523.25, 2)  
d_tone = generate_note(587.33, 2)  
e_tone = generate_note(659.25, 2)  
f_tone = generate_note(698.46, 2)  
g_tone = generate_note(783.99, 2)  
a_prime_tone = generate_note(880.00, 2)
```

### 1.4 Problem 4

- Implement `SoundWave.__add__()`.
- Generate a three-second A minor chord (A, C, and E) and embed it in the first cell below.
- Implement `SoundWave.__rshift__()`.
- Generate the arpeggio  $A \rightarrow C \rightarrow E$ , where each tone lasts one second, and embed it in the second cell below.

```
[33]: # Fundamental frequencies  
# Source: https://www1.swarthmore.edu/NatSci/ceverba1/Class/e5_2006/  
↪MusicalScales.html  
a = generate_note(440, 3) # A4  
c = generate_note(261.63, 3) # C4  
e = generate_note(329.63, 3) # E4
```

```

chord = a + c + e
chord.export("am_chord.wav", force=True)
display(Audio("am_chord.wav", autoplay=False, element_id="am-chord"))

```

Exported with scaling: am\_chord.wav

<IPython.lib.display.Audio object>

```

[34]: a1 = generate_note(440, 1)
      c1 = generate_note(261.63, 1)
      e1 = generate_note(329.63, 1)

      arpeggio = a1 >> c1 >> e1
      arpeggio.export("am_arpeggio.wav", force=True)
      display(Audio("am_arpeggio.wav", autoplay=False, element_id="am-arpeggio"))

```

Exported with scaling: am\_arpeggio.wav

<IPython.lib.display.Audio object>

## 1.5 Problem 5

- Implement `simple_dft()` with the formula  $c = F_n f$ , where  $F_n$  is the  $n$ -dimensional DFT matrix.
- In the cell below, use `np.allclose()` to check that `simple_dft()` and `scipy.fftpack.fft()` give the same result (after scaling).

```

[35]: def simple_dft(samples):
      """Compute the DFT of an array of samples.

      Parameters:
          samples ((n,) ndarray): an array of samples.

      Returns:
          ((n,) ndarray): The DFT of the given array.
      """
      n = len(samples)
      j, k = np.meshgrid(np.arange(n), np.arange(n)) # shape (n, n)
      W = np.exp(-2j * np.pi * j * k / n)           # DFT matrix F_n
      return W @ samples

```

```

[36]: samples = np.random.rand(64)
      assert np.allclose(simple_dft(samples), fft(samples))
      # Should return True if `simple_dft()` and `scipy.fftpack.fft()` give the same
      ↪ result (after scaling).

```

## 1.6 Problem 6

- Implement `simple_fft()`.

- In the cell below:
  - Generate an array of 8192 random samples and take its DFT using `simple_dft()`, `simple_fft()`, and `scipy.fftpack.fft()`.
  - Print the runtimes of each computation.
  - Use `np.allclose()` to check that `simple_fft()` and `scipy.fftpack.fft()` give the same result (after scaling).

```
[37]: def simple_fft(samples, threshold=1):
      """Compute the DFT using the FFT algorithm.

      Parameters:
          samples ((n,) ndarray): an array of samples.
          threshold (int): when a subarray of samples has fewer
                           elements than this integer, use simple_dft() to
                           compute the DFT of that subarray.

      Returns:
          ((n,) ndarray): The DFT of the given array.
      """
      n = len(samples)
      if n <= threshold:
          return simple_dft(samples)

      if n % 2 != 0:
          raise ValueError("Input size must be a power of 2")

      even = simple_fft(samples[::2], threshold)
      odd = simple_fft(samples[1::2], threshold)

      factor = np.exp(-2j * np.pi * np.arange(n) / n)
      return np.concatenate([even + factor[:n // 2] * odd,
                             even - factor[:n // 2] * odd])
```

```
[38]: import time

      # Generate random signal of length 8192
      samples = np.random.rand(8192)

      # Time simple_dft
      start = time.time()
      dft_result = simple_dft(samples)
      dft_time = time.time() - start

      # Time simple_fft
      start = time.time()
      fft_result = simple_fft(samples, threshold=32) # threshold chosen for speed
      fft_time = time.time() - start
```

```

# Time scipy.fft
start = time.time()
scipy_result = fft(samples)
scipy_time = time.time() - start

# Output runtimes
print(f"simple_dft: {dft_time:.4f} s")
print(f"simple_fft: {fft_time:.4f} s")
print(f"scipy.fft: {scipy_time:.4f} s")

# Check correctness
print("simple_fft  scipy.fft:", np.allclose(fft_result, scipy_result))

```

```

simple_dft: 5.3490 s
simple_fft: 0.0130 s
scipy.fft: 0.0002 s
simple_fft  scipy.fft: True

```

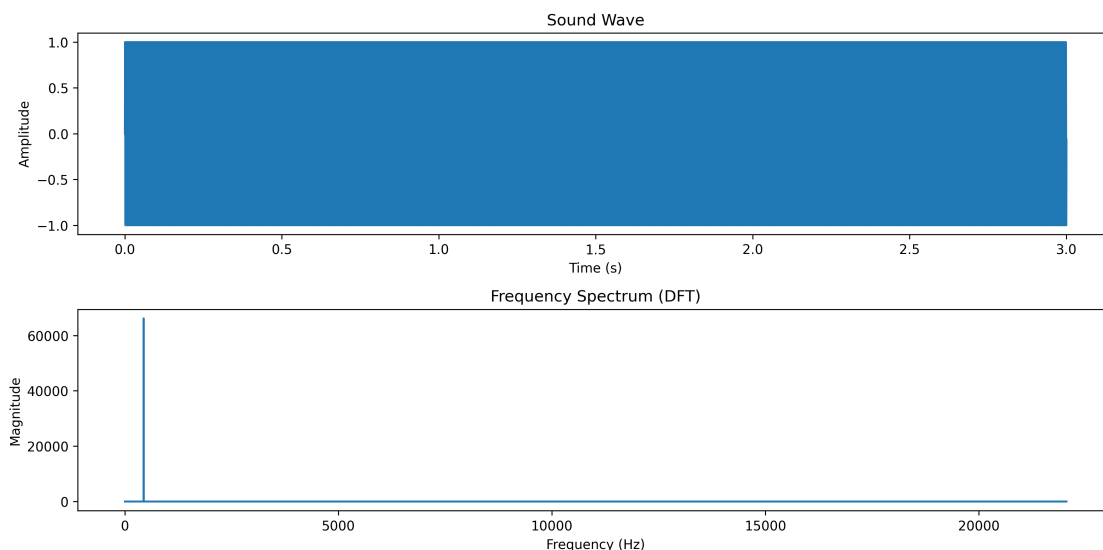
## 1.7 Problem 7

- Modify `SoundWave.plot()` so that it accepts a boolean. When the boolean is `True`, take the DFT of the stored samples and plot (in a new subplot) the frequencies present on the  $x$ -axis and the magnitudes of those frequencies on the  $y$ -axis. Only the display the first half of the plot, and adjust the  $x$ -axis so that it correctly shows the frequencies in Hertz.
- Display the plot of the DFT of the A tone from Problem 4 the first cell below.
- Display the plot of the DFT of the A minor chord from Problem 4 in the second cell below.

```

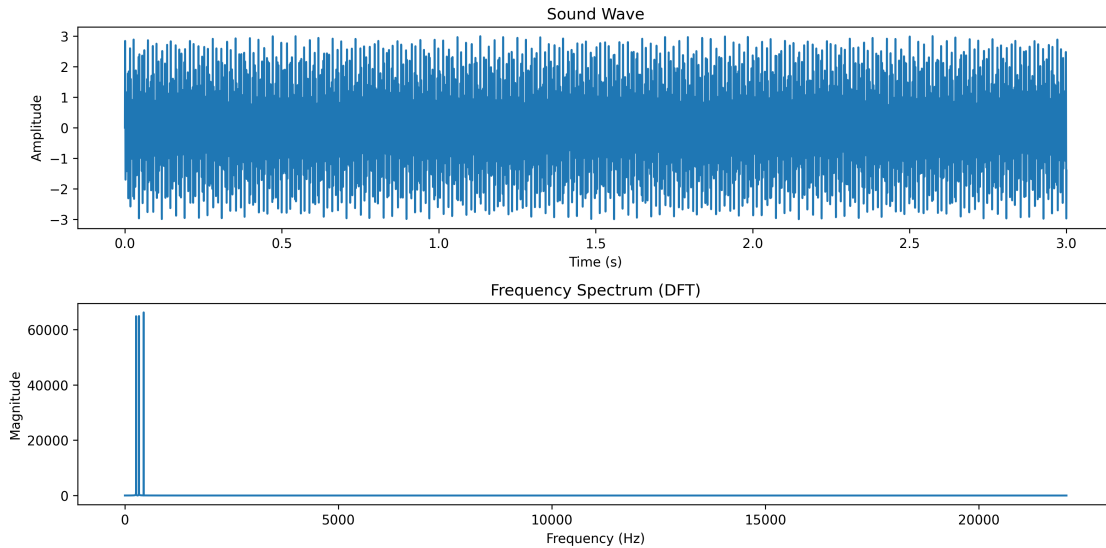
[39]: a.plot(show_dft=True)
plt.show()

```





```
[40]: chord.plot(show_dft=True)
plt.show()
```



## 1.8 Problem 8

Use the DFT to determine the individual notes that are present in `mystery_chord.wav`.

```
[41]: note_freqs = {
    "A": 440.00,
    "B": 493.88,
    "C": 523.25,
    "D": 587.33,
    "E": 659.25,
    "F": 698.46,
    "G": 783.99,
    "A ": 880.00
}

rate, samples = wavfile.read("mystery_chord.wav")
mystery_chord = SoundWave(rate, samples)

mystery_chord.plot(show_dft=True)
plt.show()

n = len(mystery_chord.samples)

freqs = mystery_chord.freq
```

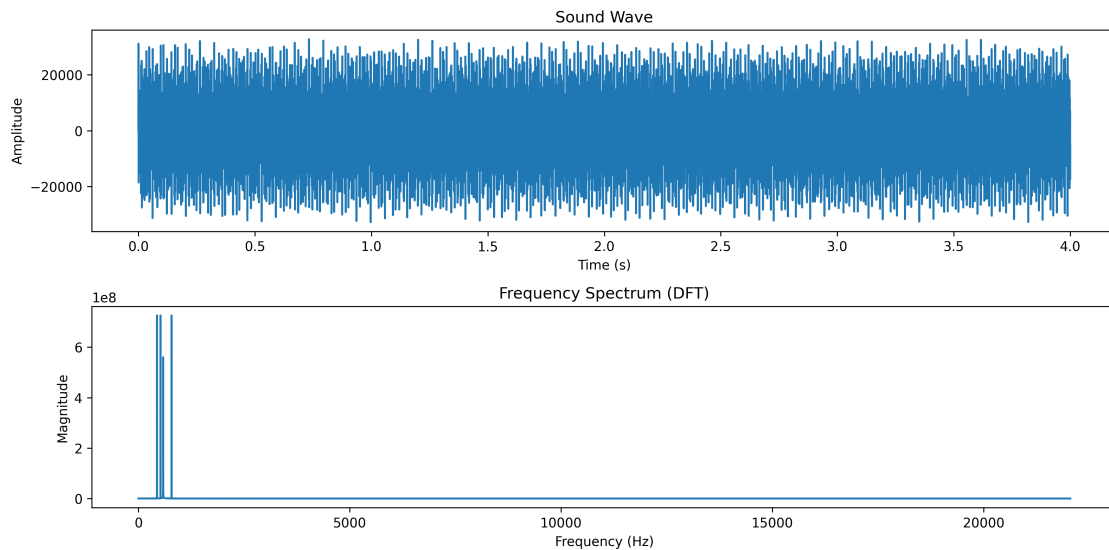
```

magnitudes = np.abs(np.fft.rfft(mystery_chord.samples))
magnitudes[0] = 0 # Ignore DC component
freqs = np.fft.rfftfreq(n, d=1/rate)
top_indices = np.argsort(magnitudes)[-10:][::-1]
top_freqs = freqs[top_indices]
top_freqs = np.unique(np.round(top_freqs, 1))[:5] # Round and deduplicate

detected_notes = []
for f in top_freqs:
    closest_note = min(note_freqs, key=lambda note: abs(note_freqs[note] - f))
    detected_notes.append((closest_note, f))
print("Detected notes and frequencies:")
for note, freq in detected_notes:
    print(f"{note}: {freq:.2f} Hz")

top_indices = np.argsort(magnitudes)[-10:][::-1]
top_freqs = freqs[top_indices]

```



Detected notes and frequencies:

A: 440.00 Hz  
C: 523.20 Hz  
D: 586.80 Hz  
D: 587.00 Hz  
D: 587.20 Hz